

Keeping Student Data Safe with AI Tools

A one-page checklist for K-12 staff using Claude, ChatGPT, Gemini or any other AI assistant.

The rule underneath everything below: **anything the AI reads goes to the vendor's servers.** That includes what you type, files you upload, and anything an AI agent opens on your computer. Keep student data out of what the AI reads, and most of the risk goes away.

1. Set up the account before any student data touches it

- **Use a district account with a signed data privacy agreement.** Treat the AI vendor like any other edtech vendor: an education or business plan, under contract, owned by the district. Don't use personal accounts for student records.
- **Turn off model training.** On personal plans it's a setting (in Claude: **Settings › Privacy › Help improve Claude**). Education and business plans generally don't train on your data by default, but confirm it in the agreement.
- **Don't rate student-data chats.** Thumbs up or down sends that conversation to the vendor as feedback, which can be used even with training turned off.
- **Delete chats you're done with.** If a process needs to be repeated, write the steps in a README file, not a chat history.

2. Send less

- **Export only the fields the task needs.** Need averages? Export student number and scores, not names, birth dates, guardians or emails.
- **Use codes instead of names.** Student numbers or made-up codes (S-014). The key that maps codes to names stays on your computer.
- **Watch for indirect identifiers.** In a small district, initials plus grade plus a 504 flag can identify a child. Drop columns you don't need.
- **Read-only access to your SIS.** AI tools should never write back to your system of record.
- **Use fake data for demos, testing and training.** Never put real student names on a projector or in a shared file.

3. Ask for a script, not a look

If your AI tool can run code on your own computer (an agent with a connected folder, such as Claude's Cowork), have it write the code and let your computer do the work.

Don't:

"Here's a spreadsheet of student names and scores. Find each student's average."

Every name and score goes to the vendor's servers and stays in your chat history.

Do:

"Don't open grades.xlsx yourself. Write a script that runs on my computer to:

1. make a copy with each name swapped for a code, and save the key as key.csv
2. average the scores in the coded copy
3. put the names back from the key and save grades_averaged.xlsx

If you need to check the data, look only at the coded copy. Never open key.csv."

The AI sees the script and a row count, at most codes and scores, never a name.

A prompt is a promise, not a lock. The AI follows instructions, but nothing physically stops it from opening a file. The safest data is data you never export.

Quick checklist

- ☐ District account with a signed data privacy agreement
- ☐ Model training turned off
- ☐ No thumbs up/down on student-data chats
- ☐ Old chats deleted; processes saved in a README
- ☐ Only the fields the task needs; codes instead of names
- ☐ Read-only access to the SIS
- ☐ Scripts run on your computer; the AI doesn't read raw student files
- ☐ Fake data for demos and testing
- ☐ Student files off personal cloud storage
- ☐ Data and secrets excluded from code repositories
- ☐ Drafts, not sends
- ☐ Reports labeled for administrators only

Cory Houghton, Technology Director, Earlham Community School District · ITEC 2026. This is practical guidance, not legal advice. Check your district's policies, FERPA and your state's student data privacy law.